

# Lab 02: Document Clustering with Hadoop MapReduce

## Lab Instructors:

**Đỗ Trọng Lễ**  
dtle@selab.hcmus.edu.vn

**Bùi Huỳnh Trung Nam**  
huynhtrungnam2001@gmail.com

## Abstract

This lab work will guide you through applying the **k-means clustering algorithm** on document datasets using Hadoop MapReduce. It consists of two parts: **data preprocessing** and **algorithm implementation**.

*This lab draws inspiration from the University of Washington's CSE547: Machine Learning for Big Data course.*

## Background knowledge

### TF-IDF (Term Frequency-Inverse Document Frequency)

**TF-IDF** is a numerical statistic that reflects the importance of a term within a document relative to a collection of documents, often used in information retrieval and text mining. It includes two components: **Term Frequency (TF)** and **Inverse Document Frequency (IDF)**.

#### Term Frequency (TF):

- Measures how often a term occurs in a document.
- Calculated as the ratio of the number of times a term appears in a document to the total number of terms in that document.
- Emphasizes the significance of a term within a specific document.

#### Inverse Document Frequency (IDF):

- Measures the rarity of a term across all documents in a collection.
- Calculated as the logarithm of the ratio of the total number of documents to the number of documents containing the term.
- Highlights terms that are unique or infrequent in the entire document corpus.

The product of TF and IDF results in a weighted score for each term, providing a measure of its importance in a specific document and across the entire document collection.

### Document Clustering

**Document clustering**, also known as **text clustering** or text categorization, is a natural language processing (NLP) technique that groups similar documents together based on their content. The goal is to discover inherent structures in a collection of documents, facilitating organization, search, and analysis. Clustering enables the identification of relationships and patterns within unstructured textual data.

## Part 1: Data Preprocessing (5 pts)

In this part, we will process textual data and convert each text document into a vector representation suitable for clustering using TF-IDF feature.

### Tasks:

#### 1.1 Text Cleaning and Term Frequency (1 pt):

The dataset we will consider comes from the BBC (<http://mlg.ucd.ie/datasets/bbc.html>).

- **Input:** Path of the directory containing the dataset. Download raw text files from the dataset page.
- **Output:** An **MTX file (Matrix Market format)** with three columns:
  - *termid*: Unique identifier for each term (word) in the vocabulary.
  - *docid*: Unique identifier for each document.
  - *frequency*: Frequency of each term in a specific document.

To understand the intended output file format, please refer to the .mtx file in the preprocessed dataset from the dataset page.

- **Steps:**
  1. Read the text data.
  2. Remove stop words (common words like "the," "a," "an"). (Download the [stop word list](#) here)
  3. Tokenize the text (split into individual words).
  4. Calculate the frequency of each term within each document.
  5. Write the output to ***task\_1\_1.mtx*** with *termid*, *docid*, and *frequency*. Get *termid* from *.terms* file and *docid* from *.docs* file. The id corresponds to the line number in the file.

#### 1.2 Low-Frequency Term Elimination (1 pt):

- **Input:** MTX file generated from Task 1.1.
- **Output:** MTX file with terms having a frequency < 3 removed.
- **Steps:**
  1. Read the MTX file.
  2. Filter terms where the frequency is less than 3 across all documents.
  3. Write the updated MTX file to ***task\_1\_2.mtx*** excluding low-frequency terms.

#### 1.3 Top 10 Most Frequent Words (1 pt):

- **Input:** MTX file from Task 1.2.
- **Output:** List of the top 10 most frequently occurring terms along with their frequencies.
- **Steps:**

1. Read the MTX file.
2. Calculate the total frequency of each term **across all documents**.
3. Sort the terms in descending order based on their total frequency.
4. Output the top 10 most frequent terms and their corresponding frequencies to **task\_1\_3.txt**.

#### 1.4 TF-IDF (1 pt):

**Term Frequency (TF):** Measures how often a term appears in a document. It is calculated as the ratio of the number of occurrences of a term to the total number of terms in the document.  $f(t, d)$  is the number of times term  $t$  appearing in document  $d$ .

$$tf(t, d) = \frac{f(t, d)}{\sum\{f(w, d): (w \in d)\}}$$

**Inverse Document Frequency (IDF):** Measures the importance of a term across a collection of documents. It is calculated as the logarithm of the ratio of the total number of documents to the number of documents containing the term.  $idf(t, D)$  is a measure of how unique or important the term  $t$  is across the entire document collection  $D$ .

$$idf(t, D) = \log \frac{|D|}{|\{d \in D: t \in d\}|}$$

**TF-IDF Score:** Combines both TF and IDF to give a weighted score that highlights terms that are important in a document but not too common across all documents.

$$tfidf(t, d, D) = tf(t, d) \times idf(t, D)$$

- **Input:** MTX file from Task 1.2.
- **Output:** Convert the *term-doc-frequency* matrix into a *term-doc-tfidf* matrix.
- **Steps:**
  1. Read the MTX file.
  2. Calculate the TF-IDF score for each term-document pair.
  3. Write the final MTX file to **task\_1\_4.mtx**.

#### 1.5 Highest average tfidf (1 pt)

For each term  $t$ , take the average tfidf over each class  $C_i = \{\text{documents in class } i\}$

$$avg\_tfidf(t, C_i, D) = \frac{1}{|C_i|} \sum_{d \in C_i} tfidf(t, d, D)$$

For each class  $C_i$ , write the 5 terms with the highest average tfidf for the class to **task\_1\_5.txt** (e.g *Tech: spywar:0.69, aol:0.58, : : : Business: : : :*).

## Part 2: K-Means Algorithm (4 pts)

In this part, you will implement the K-Means algorithm to cluster your preprocessed data.

### Tasks:

#### 2.1 K-Means on 2D Data (1.5 pt):

- **Input:** A sample file containing 2D data points (one point per line, separated by spaces).
- **Output:** Cluster assignments for each data point.
- **Steps:**
  1. Define the desired number of clusters (k).
  2. Implement the K-Means algorithm:
    - Initialize k centroids randomly.
    - Assign each data point to the closest centroid based on Euclidean distance.
    - Recompute centroids as the mean of points assigned to each cluster.
    - Repeat steps 2.2 and 2.3 until convergence (centroids no longer change significantly).
  3. Run the K-Means with K=3 for 20 iterations. Output the clusters centers to *task\_2\_1.clusters* and cluster assignment for each data point to *task\_2\_1.classes* in text format.

#### 2.2 K-Means on Preprocessed Data (1.5 pt):

- **Input:** TF-IDF MTX file from Task 1.4. For convenience, you can convert data to tfidf.txt: Each row is in the form of *docid|termid1:tfidf1,termid2:tfidf2,: : .*
- **Output:** Cluster assignments for each document represented in the MTX file and mean of each clusters.
- **Steps:**
  1. Define the desired number of clusters (k).
  2. Implement the K-Means algorithm (similar to Task 2.1) but calculate distances using a suitable distance metric for TF-IDF vectors (e.g., cosine similarity).
  3. Run k-means with K = 5 for 10 iterations. Output the clusters centers to *task\_2\_2.clusters* and cluster assignment for each data point to *task\_2\_2.classes* in text format. For each iteration, report mean (top 10 words in tfidf) of each clusters to *task\_2\_2.txt* and objective function value to *task\_2\_2.loss* in text format.
- **FAQ:**
  1. **What is the objective loss function value?**

In K-Means clustering, the objective function, often referred to as the **within-cluster sum of squares (WCSS)**, is used to evaluate the quality of the clusters formed. The goal of K-Means is to minimize this objective function value at each step. **The WCSS is defined as the sum of the squared distances between each point and the centroid of its assigned cluster.**

Here's the mathematical expression for the WCSS objective function:

$$\text{WCSS} = \sum_{i=1}^k \sum_{x \in C_i} \|x - \mu_i\|^2$$

Where:

- $k$  is the number of clusters.
- $C_i$  is the set of points in the  $i^{th}$  cluster.
- $x$  is a data point in cluster  $C_i$ .
- $\mu_i$  is the centroid of the  $i^{th}$  cluster.
- $\|x - \mu_i\|^2$  is the squared Euclidean distance between a data point  $x$  and the centroid  $\mu_i$ .

## 2. What does "report mean (top 10 words in tfidf) of each clusters to task\_2\_2.txt" mean?

After clustering, you can look at the centroids of each cluster to determine the most important terms (words) for each cluster. The centroids are essentially vectors that contain the average TF-IDF score for each term in the cluster. For each cluster centroid, sort the terms by their TF-IDF score in descending order. The top terms will be the ones with the highest scores, indicating their importance in the cluster.

### 2.3 Scalable K-Means++ Initialization (1 pt):

- **Read the article** "[Scalable K-Means++](#)" to understand the K-Means|| algorithm for improved centroid initialization.
- Implement K-Means|| to initialize centroids for your K-Means algorithm.
- Run k-means with  $K = 5$  for 10 iterations. For each iteration, report mean (top 10 words in tfidf) of each clusters to *task\_2\_3.txt* and objective function value to *task\_2\_3.loss* in text format.

## Report (1.5 pt)

The report for this lab on Document Clustering with Hadoop MapReduce should address the following:

- **Data Description:** Briefly describe the text data used for the exercise.
- **MapReduce Job Implementation:**
  - Describe the overall design of your MapReduce job, including the number of mappers and reducers used.
  - For each task:
    - Explain the logic implemented in the Map function.

- Explain the logic implemented in the Reduce function.
  - Describe the format of the key-value pairs emitted at each stage.
- **Results:**
  - Include a sample of the output from each MapReduce task.
  - Briefly discuss any challenges faced during the implementation and how you addressed them.
- **Conclusion:**
  - Summarize the key learnings from this lab.

### **Additional Considerations:**

- Include any relevant code snippets for the Map and Reduce functions (you can use pseudocode if actual code is too lengthy).
- Mention any assumptions made during the implementation.
- Maintain a clear and organized structure in your report, making it easy for the reader to follow the steps involved.

### **References.**

- Properly acknowledge any reference code utilized in your work within this section; failure to do so may be construed as **academic dishonesty**.